

Module Introduction: Operating System

- *Relevance of Operating System*

Q. You must have interacted with the operating systems when working on your laptop or mobile phone like Windows, Mac OS, Android, iOS, etc.

Think of a common scenario of copying a file and pasting the contents to another file in your computer system. Do you think operating systems play any role in performing this task? If yes, explain how operating systems do that.

Moreover, use this opportunity to think of more situations in which operating systems play a key role.

Yes, operating systems play a crucial role in performing tasks like copying and pasting files. Let's take the scenario you mentioned of copying the contents of one file and pasting them into another file.

When you select a file and choose the "Copy" option, the operating system takes care of several tasks behind the scenes:

1. **File Management:** The operating system manages the file system, which includes organizing and keeping track of files on your storage device. It locates the source file and retrieves its contents.
2. **Clipboard Management:** The operating system provides a clipboard where the copied content is temporarily stored. The clipboard acts as a buffer for transferring data between applications. When you choose the "Copy" option, the operating system places the selected file's content into the clipboard.
3. **User Interface Integration:** The operating system ensures that the clipboard is accessible to applications through a standardized user interface. This allows you to interact with the clipboard using common commands like "Copy" and "Paste" across different applications.
4. **Content Transfer:** When you choose the "Paste" option, the operating system retrieves the content from the clipboard and makes it available to the target application. The operating system handles the transfer of the content, allowing the target file to receive the copied data.

Operating systems also play key roles in various other scenarios:

1. **Process and Memory Management:** Operating systems manage the execution of processes and allocate system resources like memory to ensure efficient utilization of hardware.

2. **Device Drivers:** Operating systems provide drivers that allow hardware devices (e.g., printers, scanners, graphics cards) to communicate with the computer. The operating system facilitates the installation, configuration, and management of these drivers.

3. **Networking:** Operating systems provide networking capabilities, allowing devices to connect and communicate with each other. They manage network protocols, handle data transmission, and provide security features.

4. **Security:** Operating systems implement security measures such as user authentication, file permissions, and encryption to protect data and system resources from unauthorized access.

5. **File System Management:** Operating systems handle file organization, access control, and storage management. They ensure that files are stored efficiently, provide mechanisms for searching and retrieving files, and enforce file-level security.

6. **User Interface:** Operating systems provide a graphical user interface (GUI) or command-line interface (CLI) through which users interact with the computer system. They manage windows, menus, icons, and other visual elements to facilitate user interaction.

These are just a few examples of how operating systems play a vital role in various aspects of computing, from basic file management to advanced system operations.

Q. Look at the following tasks done by the OS:

1. The distribution of memory among different applications

2. Different applications running on the computer system

These examples were explained in the lecture video. In both of these examples, OS was doing abstraction as well as arbitration.

How are these two examples different from each other? Also, can you justify the functions of the OS in these examples?

In the given examples, both tasks involve the operating system performing abstraction and arbitration, but they differ in terms of the specific functions and purposes they serve.

1. Distribution of Memory Among Different Applications:

In this example, the operating system performs abstraction by presenting a virtualized view of the memory to each application. It hides the physical complexities of the memory system and provides a logical address space for each application to use. Each application is unaware of the actual physical addresses and shares the illusion of having exclusive access to the entire memory.

The function of the operating system in this example is memory management. It ensures that memory resources are allocated efficiently among different applications, preventing them from interfering with each other's memory space. The operating system sets up memory partitions or

uses techniques like virtual memory to allocate and manage memory for each application, enforcing memory protection and facilitating secure and isolated execution of applications.

2. Different Applications Running on the Computer System:

In this example, the operating system performs abstraction by providing a consistent and standardized interface for applications to run on the computer system. It abstracts away the underlying hardware complexities, allowing applications to run on different hardware configurations without needing to be rewritten or recompiled for each specific system.

The function of the operating system in this example is process management and system scheduling. It manages the execution of multiple applications or processes concurrently, ensuring fair and efficient utilization of the CPU and other system resources. The operating system schedules processes based on priority, time-sharing, or other scheduling algorithms, providing each application with its fair share of CPU time and ensuring a responsive and smooth user experience.

In both cases, the operating system also performs arbitration, which involves mediating access to shared resources and resolving conflicts. For example, in memory distribution, the operating system arbitrates among applications to allocate memory resources based on their needs and priorities. In the case of running different applications, the operating system arbitrates CPU time and system resources to ensure fair and efficient execution of multiple applications.

Overall, the operating system functions as a memory manager and process scheduler in the distribution of memory among applications and the running of different applications, respectively. It abstracts away hardware complexities, provides a standardized interface, and arbitrates resource allocation and scheduling to facilitate efficient and secure execution of applications on the computer system.

● ***Abstraction and Arbitration***

Abstraction and arbitration are two important concepts in the context of operating systems. Let's define and understand each concept:

Abstraction: Abstraction in operating systems refers to the process of hiding the complex details of the underlying hardware or software and providing a simplified and standardized interface to users or applications. It allows users and applications to interact with the system without needing to understand the intricate technical aspects.

Abstraction helps in simplifying the programming and usage of computer systems. By providing a higher-level view and hiding unnecessary complexities, it makes it easier for users and applications to work with the system. Examples of abstractions provided by the operating system include file systems, memory management, device drivers, and network protocols.

Arbitration: Arbitration in operating systems refers to the process of mediating and controlling access to shared resources or resolving conflicts that arise when multiple entities compete for the same resource. The operating system acts as a referee, ensuring fair and controlled access to shared resources.

Arbitration is necessary because multiple processes or applications running concurrently on a system may need access to shared resources, such as the CPU, memory, or I/O devices. The operating system manages the allocation and scheduling of these resources, ensuring that conflicts are resolved and resources are efficiently utilized.

The operating system uses various algorithms and techniques for resource arbitration, such as priority-based scheduling, time-sharing, or queuing mechanisms. These mechanisms help determine the order and manner in which processes or applications are granted access to resources, ensuring fairness, and efficiency, and preventing issues like starvation or resource deadlock.

Abstraction and arbitration are closely related and often go hand in hand. Abstraction simplifies the usage of resources by providing a higher-level interface, while arbitration ensures that resources are shared and accessed in a controlled and fair manner.

In summary, abstraction hides complexities and provides a simplified interface, while arbitration mediates access to shared resources. Both concepts are essential for the smooth functioning of operating systems and for facilitating efficient and effective utilization of computer resources.

kernel

The kernel is a core component of an operating system. It is the central part of the operating system that provides essential services and manages the system's resources. The kernel acts as a bridge between applications and the underlying hardware of a computer system.

Here are some key characteristics and functions of the kernel:

- **Resource Management:** The kernel manages system resources such as CPU, memory, disk space, and I/O devices. It allocates and schedules these resources among different applications or processes running on the system.
- **Process Management:** The kernel oversees the creation, execution, and termination of processes. It manages the execution context of processes, including their memory space, CPU time, and I/O operations.

- **Memory Management:** The kernel is responsible for memory allocation and management. It assigns memory to processes, keeps track of allocated memory, and ensures memory protection and security.
- **Device Drivers:** The kernel provides device drivers that enable communication between the operating system and hardware devices. These drivers handle device-specific operations and facilitate data transfer between applications and devices.
- **File System Management:** The kernel manages file systems, which include organizing and accessing files on storage devices. It provides services for file creation, deletion, reading, and writing, and enforces file-level permissions and security.
- **System Calls:** The kernel exposes a set of system calls, which are interfaces that allow applications to request services from the operating system. System calls provide a way for applications to interact with the kernel and access its functionalities.
- **Kernel Security:** The kernel implements security measures to protect the system and its resources. It enforces access control, manages user authentication, and provides mechanisms for secure communication and data protection.
- **Kernel-level Interrupt Handling:** The kernel handles interrupts generated by hardware devices or exceptional events. It responds to interrupts promptly, ensuring timely processing and appropriate actions.

The kernel is typically loaded into memory during system startup and remains in memory throughout the system's operation. It operates in privileged mode, allowing it to access and control critical system resources. The kernel interacts with applications, device drivers, and other components of the operating system to provide a cohesive and functional computing environment.

Different operating systems have different kernel designs, such as monolithic kernels, microkernels, or hybrid kernels. The specific design and implementation of the kernel depend on the goals, requirements, and architecture of the operating system.

There are different types of kernels used in operating systems, each with its own design and characteristics. The main types of kernels are:

1. Monolithic Kernel:

A monolithic kernel is a single, large software module that provides all core operating system functionalities as a single unit. It runs in kernel mode and handles process management, memory management, file system management, device drivers, and other system services. Monolithic kernels are efficient and provide low overhead since they allow direct access to

hardware resources. However, their large size makes them more susceptible to bugs and can lead to less modular and flexible operating systems.

2. Microkernel:

A microkernel is designed to be minimalistic, with only essential functionalities implemented in the kernel. It provides a minimal set of services, such as process scheduling, interprocess communication (IPC), and basic memory management. Other functionalities, such as file systems and device drivers, are implemented as user-level processes or modules running outside the kernel. Microkernels aim to provide a more modular and extensible operating system architecture, allowing easier maintenance, scalability, and fault tolerance. However, the communication overhead between user-level processes can impact performance compared to monolithic kernels.

3. Hybrid Kernel:

A hybrid kernel combines elements of both monolithic and microkernel designs. It includes a relatively small kernel that provides essential services such as process scheduling, memory management, and IPC. However, it also includes some additional services traditionally found in monolithic kernels, like device drivers and file systems, which run in kernel space. Hybrid kernels aim to strike a balance between the efficiency of monolithic kernels and the modularity of microkernels, offering both performance and flexibility.

4. Exokernel:

An exokernel takes the idea of minimalism to the extreme by providing only the barest abstractions to applications. It exposes hardware resources directly to applications, allowing them to have fine-grained control over resource management. Exokernels essentially act as a thin layer between the hardware and the application, providing only basic protection and resource multiplexing. This design allows for maximum flexibility and customization but requires applications to handle most of the operating system functionality themselves.

These are the main types of kernels, each with its own trade-offs and design principles. The choice of kernel type depends on factors such as performance requirements, system architecture, development goals, and the desired balance between efficiency and modularity. Different operating systems may employ different kernel types to suit their specific needs.

Why do you think the Hybrid kernel has better performance than Microkernel and small size in comparison to the Monolithic kernel?

The hybrid kernel offers a better performance compared to a microkernel and a smaller size compared to a monolithic kernel due to the following reasons:

1. *Reduced Communication Overhead:* In a hybrid kernel, essential services like process scheduling and memory management are handled within the kernel, similar to a monolithic kernel. This reduces the need for frequent communication between user-level processes or

modules, which can introduce overhead in a microkernel design. By minimizing communication overhead, the hybrid kernel can achieve better performance.

2. *Direct Access to Hardware*: The hybrid kernel allows certain critical services, such as device drivers and file systems, to run in kernel space, similar to a monolithic kernel. This enables these services to have direct access to hardware resources, resulting in efficient and optimized handling of I/O operations. By avoiding the need to go through user-level processes or modules, the hybrid kernel can provide better performance compared to a microkernel.

3. *Flexibility and Modularity*: While a monolithic kernel may have a larger size due to incorporating all functionalities into a single module, a hybrid kernel offers greater modularity. By separating essential services from additional functionalities, the hybrid kernel allows for better organization and easier maintenance. This modular design also enables better scalability, as additional features can be added without impacting the core kernel components. As a result, the hybrid kernel can have a smaller size compared to a monolithic kernel while still providing comparable performance.

Overall, the hybrid kernel strikes a balance between the performance advantages of a monolithic kernel and the modularity benefits of a microkernel. It retains critical services within the kernel for improved performance, while allowing additional functionalities to be implemented as separate modules for flexibility and easier maintenance.

System call and it's architecture

A system call is a mechanism provided by the operating system that allows user-level applications to request services or functionality from the kernel. It serves as an interface between user-space and kernel-space, enabling applications to access privileged operations and interact with the underlying operating system.

When an application wants to perform a privileged operation or access a resource managed by the operating system, it initiates a system call. The application makes a request to the kernel by invoking a specific system call function provided by the operating system's API (Application Programming Interface). The system call function typically takes parameters specifying the requested operation and returns a result or status.

The architecture of a system call involves several key components:

1. *User-Space Application*: The user-space application is the program or process running in the user mode, outside the kernel. It initiates a system call when it needs to access privileged operations or resources provided by the operating system.

2. **System Call Interface:** The system call interface is a collection of functions, typically provided as a library or set of APIs, that allow applications to make system calls. The interface provides a set of well-defined functions, each corresponding to a specific system call operation. Examples of system call interfaces include POSIX API for Unix-like systems and WinAPI for Windows.

3. **System Call Handler:** The system call handler is a part of the operating system's kernel. It is responsible for receiving and processing system call requests from user-space applications. When a system call is invoked, the system call handler is activated to handle the request.

4. **Transition to Kernel Mode:** To perform privileged operations, the user-space application needs to transition from user mode to kernel mode. This transition is initiated when a system call is made, triggering a switch from user-space to kernel-space.

5. **System Call Execution:** Once the system call request reaches the kernel, the system call handler identifies the requested operation based on the provided parameters. It verifies the validity of the request, checks permissions, and performs the requested operation on behalf of the user-space application.

6. **Return to User-Space:** After executing the system call, the kernel returns the result or status back to the user-space application. The system call handler copies the result to a designated memory location accessible by the application.

The architecture of the system call allows user-level applications to safely and efficiently interact with the kernel and utilize the privileged services and resources provided by the operating system. It provides a controlled and secure interface for user-space applications to access and utilize system-level functionalities.